

Cloud Computing Google App Engine



“Run your Apps, Host your Sites, Store your Data”

DI Manfred Pamsl
DI(FH) Mathias Knoll, MSc

Overview

- Platform as a Service
 - Web applications are deployed to the Google platform
 - Integration with Google accounts through APIs
- Programming Language Support
 - Java, Python, Go, PHP
- Flexible, scalable application storage
 - Relational database, cloud storage
- Developer tools
 - Eclipse plug-in, web development tools, ...

The Sandbox

- Applications run in an environment with limited access to the operating system
 - Isolates applications in an environment that is independent of the hardware, OS and physical location of the web server
 - Access to the applications only via HTTP(S)
 - Applications may only access files uploaded by the application code itself, there is no general file system access
 - Application code is only executed in response to a web request, a queued task or scheduled task - response data must be returned within 60s
- Restrictions allow Google to distribute the requests and applications across their infrastructure

Generally Available Features

- Data storage
- Communication
- Process management
- Computation
- App configuration and management

Data Storage

- Schema-less object datastore
- Accessible through SQL like language GQL
- Supports Binary Large Objects (BLOBs)
- Distributed in-memory data cache (“memcache”)
- Provides programmatic access to application logs and request logs

Communications

- Persistent communication channels between web clients (via Javascript library) and the application
- Send and receive email messages using Google mail accounts
- Issue HTTP(S) request using the Google networking infrastructure
- Support for Extensible Messaging and Presence Protocol (XMPP)

Process Management & Computation

- Perform tasks outside a user request
- Configuration of regularly scheduled tasks
- “Backends” are instances of an application which have access to more computing resources
- Image conversion (photographs, graphics) API

App Configuration & Management

- Application identity services with “Oauth” (see: <http://tools.ietf.org/html/draft-ietf-oauth-v2-22>)
- Provides detection of outages and scheduled maintenance
- Application access with HTTP(S) via a custom domain instead of the default “appspot.com”
- Remote API lets you transparently access App Engine services from any application
- Namespaces API to build multi-tenant applications
- Traffic Splitting to different versions of an application based on IP address or Cookies (“A/B Testing”)
- User authentication through Google accounts or OpenID

The Java Runtime Environment

- Applications can be developed using common Java development tools and API standards
- Interaction with the environment using Java Servlet Standard, applications may use the Java Server Pages technology
- Runtime environment uses the Java 7 standard
- Many App Engine services are accessed using standard Java APIs
 - E.g. Java Data Objects (JDO), Java Persistence API (JPA), Java Mail API, java.net HTTP APIs
- Additional APIs for App Engine services not covered by Java standards

Options for Storing Data

- App Engine Datastore
 - NoSQL schemaless object data store
- Google Cloud SQL
 - Relational SQL database service based on MySQL
- Google Cloud Storage
 - Storage Service for objects and files (comparable to OpenStack Swift)

App Engine Data Store

- Distributed NoSQL data storage service
- Features a query engine and atomic transactions
- Data objects (“entities”) are identified by keys consisting of
 - “Kind”: category of the object
 - “Identifier”: name string or automatically assigned integer value
 - “Ancestor Path”: optional, link to parent entity, locates the entity in the datastore hierarchy
- Entities may have properties assigned (e.g. integer, strings, ...)
- Queries can retrieve entities of a given kind filtered and sorted by the values of properties

Application Development

- App Engine SDK includes a web server that emulates all of the App Engine services, includes a tool to upload the application to the App Engine
- Java SDK runs on any Java 7 platform
- Google Eclipse Plugin is available to create, test and upload App Engine Applications
- Web-based Administration Console is used to manage the applications on the App Engine

Exercise Steps

- Use the provided Eclipse Environment for the App Engine application development
 - Already includes the Plugin, App Engine SDK and JDK 7
- Follow the guideline to create your first sample “guestbook” application:
<https://developers.google.com/appengine/docs/java/gettingstarted/introduction>
 - Skip the the SDK installation step, it is already included in the Eclipse bundle
 - Start with “Creating a Project”
- As your homework, create an additional “point of interest” application as defined on the next slides

Point of Interest Application

- Purpose
 - Public accessible registration of points of interests, usable by any application
- Implementation requirements
 - Implement a Java application hosted on the Google AppEngine Cloud infrastructure
 - Provide ReST Web Services, that uses data in JSON format for point of interest registration and retrieval
 - you may use the gson library: <http://code.google.com/p/google-gson/> for converting Java objects to JSON format and vice versa
 - Use the App Engine schema-less object data store for data storage

Point of Interest Attributes

- “id”: unique identifier, long value, automatically generated when adding a point of interest
- “name”: Name of the point of interest, “String” value
- “latitude”: latitude position value, “double” value
- “longitude”: longitude position value
- “creator”: person creating the entry, “String” value
- “description”: short description of the point of interest, “String” value
- “category”: category of the point of interest, “String” value

JSON POI Example

```
{  
  "id": 12346789,  
  "name": "FH-Graz",  
  "latitude": 47.0693127,  
  "longitude": 15.4079899,  
  "creator": "john",  
  "description": "FH JOANNEUM Graz",  
  "category": "Fachhochschule"  
}
```


ReST API Description

- `http://<YourGoogleAppURL>/resources/poi`
 - GET: retrieve all POIs as JSON array
 - POST: Add one POI in JSON format (attribute “id” shall be ignored), returns the JSON POI with the generated “id” attribute
- `http://<YourGoogleAppURL>/resources/poi?<attribute>=<value>...`
 - GET: returns the POIs with matching attribute values (logical “and” relation)
- `http://<YourGoogleAppURL>/resources/poi/<id>`
 - GET: retrieve the POI defined by ID <id>
 - POST: Update the POI defined by ID <id>
 - DELETE: Delete the POI defines by ID <id>
- When encountering an error, an appropriate HTTP error code (4xx) and error message shall be returned

Teamwork

- Build teams of three members
- Responsibilities:
 - First member: implements the ReST API Servlet
 - Second member: implements the internal database API used by the ReST API Servlet
 - Third member: client program for basic CRUD operations (any technology you like, e.g. native Java or Python client program, Google App Engine web interface, ...)

Deliverable

- Application has to be deployed on the Google cloud
- Upload of a ZIP file containing:
 - (Eclipse) projects including the source code
 - Document describing the client program usage and access to the ReST service